

The purpose of this section is to showcase the ability to scrape and process web data using R. This tutorial draws heavily from a blog post on a great blog by two biologists and experienced R users. Suppose we want to find the number of available environmental economics journals. There are too many. Definitely. But suppose we want to find out just how many. To do this, we can visit [crossref.org](http://crossref.org), which is a citation-linking network with a list of all journals and their Digital Object Identifiers (DOIs). We will query the list from within R and then parse the returned content to list journals with certain attributes. For this, we'll need to load the following libraries:

```
library(XML)
library(RCurl)
library(stringr)
```